



Exercise 9.1: Deploy A New Service

Overview

Services (also called **microservices**) are objects which declare a policy to access a logical set of Pods. They are typically assigned with `labels` to allow persistent access to a resource, when front or back end containers are terminated and replaced.

Native applications can use the `Endpoints` API for access. Non-native applications can use a Virtual IP-based bridge to access back end pods. `ServiceTypes` Type could be:

- **ClusterIP** default - exposes on a cluster-internal IP. Only reachable within cluster
- **NodePort** Exposes node IP at a static port. A ClusterIP is also automatically created.
- **LoadBalancer** Exposes service externally using cloud providers load balancer. NodePort and ClusterIP automatically created.
- **ExternalName** Maps service to contents of `externalName` using a CNAME record.

We use services as part of decoupling such that any agent or object can be replaced without interruption to access from client to back end application.

1. Deploy two **nginx** servers using **kubectl** and a new `.yaml` file. The kind should be `Deployment` and label it with `nginx`. Create two replicas and expose port 8080. What follows is a well documented file. There is no need to include the comments when you create the file. This file can also be found among the other examples in the tarball.

```
student@cp:~$ vim nginx-one.yaml
```

YAML

nginx-one.yaml

```
1 apiVersion: apps/v1
2 # Determines YAML versioned schema.
3 kind: Deployment
4 # Describes the resource defined in this file.
5 metadata:
6   name: nginx-one
7   labels:
8     system: secondary
9 # Required string which defines object within namespace.
10  namespace: accounting
11 # Existing namespace resource will be deployed into.
12 spec:
13   selector:
14     matchLabels:
15       system: secondary
16 # Declaration of the label for the deployment to manage
17   replicas: 2
18 # How many Pods of following containers to deploy
19   template:
20     metadata:
21       labels:
22         system: secondary
23 # Some string meaningful to users, not cluster. Keys
```



```

24 # must be unique for each object. Allows for mapping
25 # to customer needs.
26   spec:
27     containers:
28 # Array of objects describing containerized application with a Pod.
29 # Referenced with shorthand spec.template.spec.containers
30     - image: nginx:1.20.1
31 # The Docker image to deploy
32     imagePullPolicy: Always
33     name: nginx
34 # Unique name for each container, use local or Docker repo image
35     ports:
36     - containerPort: 8080
37       protocol: TCP
38 # Optional resources this container may need to function.
39     nodeSelector:
40       system: secondOne
41 # One method of node affinity.

```

2. View the existing labels on the nodes in the cluster.

```
student@cp:~$ kubectl get nodes --show-labels
```

```
<output_omitted>
```

3. Run the following command and look for the errors. Assuming there is no typo, you should have gotten an error about the accounting namespace.

```
student@cp:~$ kubectl create -f nginx-one.yaml
```

```
Error from server (NotFound): error when creating
"nginx-one.yaml": namespaces "accounting" not found
```

4. Create the namespace and try to create the deployment again. There should be no errors this time.

```
student@cp:~$ kubectl create ns accounting
```

```
namespace/accounting" created
```

```
student@cp:~$ kubectl create -f nginx-one.yaml
```

```
deployment.apps/nginx-one created
```

5. View the status of the new pods. Note they do not show a Running status.

```
student@cp:~$ kubectl -n accounting get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-one-74dd9d578d-fcpmv	0/1	Pending	0	4m
nginx-one-74dd9d578d-r2d67	0/1	Pending	0	4m

6. View the node each has been assigned to (or not) and the reason, which shows under events at the end of the output.

```
student@cp:~$ kubectl -n accounting describe pod nginx-one-74dd9d578d-fcpmv
```

```
Name:          nginx-one-74dd9d578d-fcpmv
Namespace:     accounting
Node:          <none>
```

```
<output_omitted>
```

```
Events:
```

Type	Reason	Age	From	...
Warning	FailedScheduling	<unknown>	default-scheduler	

0/2 nodes are available: 2 node(s) didn't match node selector.

7. Label the secondary node. Note the value is case sensitive. Verify the labels.

```
student@cp:~$ kubectl label node worker system=secondOne
```

```
node/worker labeled
```

```
student@cp:~$ kubectl get nodes --show-labels
```

NAME	STATUS	ROLES	AGE	VERSION	LABELS
cp	Ready	control-plane	15h	v1.25.1	beta.kubernetes.io/arch=amd64, beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=cp, kubernetes.io/os=linux,node-role.kubernetes.io/control-plane=,node-role.kubernetes.io/master=, node.kubernetes.io/exclude-from-external-load-balancers=
worker	Ready	<none>	15h	v1.25.1	beta.kubernetes.io/arch=amd64, beta.kubernetes.io/os=linux,kubernetes.io/arch=amd64,kubernetes.io/hostname=worker, kubernetes.io/os=linux,system=secondOne

8. View the pods in the accounting namespace. They may still show as Pending. Depending on how long it has been since you attempted deployment the system may not have checked for the label. If the Pods show Pending after a minute delete one of the pods. They should both show as Running after a deletion. A change in state will cause the Deployment controller to check the status of both Pods.

```
student@cp:~$ kubectl -n accounting get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-one-74dd9d578d-fcpmv	1/1	Running	0	10m
nginx-one-74dd9d578d-sts5l	1/1	Running	0	3s

9. View Pods by the label we set in the YAML file. If you look back the Pods were given a label of app=nginx.

```
student@cp:~$ kubectl get pods -l system=secondary --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
accounting	nginx-one-74dd9d578d-fcpmv	1/1	Running	0	20m
accounting	nginx-one-74dd9d578d-sts5l	1/1	Running	0	9m

10. Recall that we exposed port 8080 in the YAML file. Expose the new deployment.

```
student@cp:~$ kubectl -n accounting expose deployment nginx-one
```

```
service/nginx-one exposed
```

11. View the newly exposed endpoints. Note that port 8080 has been exposed on each Pod.

```
student@cp:~$ kubectl -n accounting get ep nginx-one
```

NAME	ENDPOINTS	AGE
nginx-one	192.168.1.72:8080,192.168.1.73:8080	47s

12. Attempt to access the Pod on port 8080, then on port 80. Even though we exposed port 8080 of the container the application within has not been configured to listen on this port. The **nginx** server listens on port 80 by default. A `curl` command to that port should return the typical welcome page.

```
student@cp:~$ curl 192.168.1.72:8080
```

```
curl: (7) Failed to connect to 192.168.1.72 port 8080: Connection refused
```

```
student@cp:~$ curl 192.168.1.72:80
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<output_omitted>
```

13. Delete the deployment. Edit the YAML file to expose port 80 and create the deployment again.

```
student@cp:~$ kubectl -n accounting delete deploy nginx-one
```

```
deployment.apps "nginx-one" deleted
```

```
student@cp:~$ vim nginx-one.yaml
```

YAML

nginx-one.yaml

```
1 ....
2     ports:
3     - containerPort: 8080    #<-- Edit this line
4       protocol: TCP
5     ....
```

```
student@cp:~$ kubectl create -f nginx-one.yaml
```

```
deployment.apps/nginx-one created
```